

THE TIMELINE

VOL. I — NO. 1

WEDNESDAY, AUGUST 19, 2026

LIMA, PERU

PRICE 25¢

ENGINEERING WRITE-UPS AND LEARNING

A Single Skill Took Claude Code to a Perfect Score on the ARC-AGI-3 Benchmark

Poobesh Gowtham reports 25 of 25 games and 183 of 183 levels finished, and a top RHAE of 100.00, using Claude Code with Opus 5 and no additional harness. The method is one rule: force a falsifiable prediction before every action.

BY @PBSHGTHM

HOME TIMELINE · RUN OF 19 AUGUST 2026, 22:24

GAMES FINISHED

25 of 25

LEVELS FINISHED

183 of 183

RHAE, THE TOP SCORE

100.00 of 100

The scoreboard as posted by @pbsghthm: a general-purpose coding agent, run without a bespoke harness, clearing every game and every level on ARC-AGI-3. Full write-up at arc-skill.vercel.app.

ARC-AGI-3 might actually be a skill issue, @pbsghthm wrote on Wednesday, and then posted the scoreboard to prove it. Running Claude Code with Opus 5 — no additional harness, and no scaffolding built for the benchmark — the agent finished 25 of 25 games and 183 of 183 levels, and took the top RHAE score of 100.00 out of 100. The benchmark has been treated as a standing rebuke to general agents. The post argues the opposite: that the gap was never capability, only method. The author has published the skill itself and invites others to reproduce the run before treating the number as settled.

The key idea was simple, the author writes: force a falsifiable prediction before every action. That turns every move into an experiment, every miss into a precise correction, and lets the agent learn the game by being wrong on the record. The claim is narrow but pointed. It is not a new model, and it is not a new harness. It is one skill file, applied to a general-purpose coding agent that thousands of developers already run at their own desks. Reaction was quick and mostly technical. The striking part is not the score but how cheap the intervention was: no fine-tuning, no bespoke tooling, nothing that could not be copied into another project in an afternoon.

The result lands in a week already thick with agent tooling. On the same timeline @Benioff put Slack Code into general release, with agents from Anthropic, GitHub, Cognition and Vercel working in the same channel as their human colleagues. @ClaudeDevs shipped a Concise output style for Claude Code, set from /config or from settings.json. And @cursor_ai published a long account of rebuilding Git hosting as a database. What the week did not produce was a new frontier model. The movement was all in the layer above: the harnesses, the skills and the plumbing that decide what a model is allowed to do.

Continued on Page 4

The Digest

Everything else worth keeping from the run of 19 August.

TOOLS THAT CONNECT TO AI AGENTS

@BENIOFF

Slack Code is live. Humans and agents in the same channel, on the same work, with agents from Anthropic, GitHub, Cognition and Vercel.

@CLAUDEDEVS

Claude Code takes a Concise output style: the result first, full detail on request. Set it in /config or in settings.json.

@JASONZHOU1993

Treg, an OpenRouter for tools: 2,600 agent-friendly APIs, searchable by task, paid per call and with no subscription. Open sourced.

@CHARLIEHOLTZ

Conductor and Paper, demonstrated as one workflow rather than two separate tools.

RELEASES AND DEVELOPER TOOLS

@ROBKNIIGHT__

terminal-code puts VS Code inside the terminal: a compatible CLI that works over ssh and takes your terminal theme.

@RAUCHG

fx.sh is 6.3 MB and starts in ten microseconds — a Zig-compiled static ELF binary, or a smaller WebAssembly build.

@VICTORTAELIN

Bend2 is down from 242,000 tokens to about 80,000. The whole repository now fits in one context window.

@NUTLOPE

1kpapers summarises and visualises the year's top thousand research papers. Total cost to summarise them: four dollars.

AI-BUSINESS CONTENT

Stripe Is Buying OpenRouter

@patrickc confirmed the acquisition on the timeline, calling OpenRouter the best way to use new models and to manage several inference providers at once. His argument for the deal is that every business will soon have to manage token flows the way it already manages revenue flows.

It is money news, which this page normally leaves out. It runs here because it changes a tool many readers route their own models through. @shivamhwp posted the joke that wrote itself: a fresh pricing sheet at 3.6 per cent plus thirty cents per API request.

NEW AI MODELS

Superwhisper Opens the Weights on S1-Mini

Superwhisper has published S1-mini, the company's first open-weights language model. It runs 0.6 billion parameters and processes transcripts entirely on the user's own device, with no audio leaving the machine. It is in the app today.

@neilsuperduper said the team deliberated for a long time internally about whether to open the weights at all. They concluded that publishing was the faster route to better open models in the category, because opening their own was the best way to push things in that direction.

ROBOTICS AND HARDWARE

A One-Shot Learner for the Physical World

@GeneralistAI introduced GEN-1.5, a one-shot learner. Show the system what to do and it learns the task in a few seconds, then generalises from that single demonstration to situations it was never shown.

The company says the capability emerged from pre-training on physical data at scale, and frames it as a step towards general intelligence for the physical world. @himanshustwts, watching the demonstration, called it an absolute GPT-3 moment for robotics.

Also on the Wire

@CURSOR_AI

Twenty years of Git infrastructure, and why Origin is operated as if it were a database. cursor.com/blog/git-at-any-scale

@LINEAR

The read path of the local-first sync engine rebuilt on turbopuffer, with inverted indexes narrowing millions of sync actions.

@EMMETTSHINE

A build breakdown of littleplains.com: the models used, the retrieval architecture, and an eight-step guide.

@CNOLANMINICH

portless retires localhost:3000 and friends for good. github.com/vercel-labs/portless

@NOTRAB

orchard, a friendlier front end for Apple Container on a new Mac, in place of Docker. orchard.andon.dev

@DAVIDONDREJ1

bb, the mobile client for agents, put forward for a proper trial rather than a glance.

@ZAIDMUKADDAM

learn-inference.com, a free and unusually deep course in inference engineering.

@ASKALPHAXIV

Paperscrolling: trending research with key ideas, figures and audio explanations from alphaXiv Briefs.

@NICKADOBOS

Cerebras CS-4 clocked at 1,400 tokens a second, against roughly 50 to 70 for Sonnet 5.

@RWALK_XYZ

A free, CC0-licensed legal template library for startups: MSAs, employment agreements and the rest.

@MDP

cap-table, an open model for founders who do not yet understand their own. github.com/1984vc/cap-table

@GEMINIAPP

A free year of Google AI Pro for students in the United States, and AI Plus in 140 further countries.

THE DESK'S NOTE

This page was set from a single read of the X home timeline, run 2026-08-19_2224, which returned 127 new posts once advertisements, pagination cursors, duplicates and previously seen ids had been removed. Twenty-four were kept and printed here.

The desk skipped funding and valuation news, model hot takes, vendor drama, reactions without the artifact attached, politics, memes and personal posts. The Stripe and OpenRouter item breaks that rule on purpose, because it changes a tool many readers route their own models through, and it is marked as such.